

Position Sensitivity and Lexical Competition in Synthetic Clause Retrieval with a Small Language Model

Anonymous Authors¹

Abstract

Where an answer appears and which competing text survives selection may govern clause retrieval more clearly than prompt length in a small language model. We evaluate the frozen Qwen2.5-0.5B-Instruct Q8 artifact on 180 deterministic synthetic contract-clause probes spanning mean prompt lengths of 1989, 3986, and 7992 tokens. Moving the answer clause from the start to the end raised accuracy by 58.3 points at the shortest tier and 50.0 points at the middle tier. Keyword pruning reduced accuracy from 40.0% to 6.7% because it retained a clause whose wording matched the question but whose code was wrong. A query-entity selector excluded that lexical decoy and reached 100.0% accuracy on the same 30 cases. The matched length comparison remained inconclusive, so the practical finding is narrower: test a cheap selector against plausible lexical competitors before using it to shorten documents for this model.

1. Introduction

You paste a long contract into a chatbot and ask, “What code renews the vendor agreement?” The answer sits in one clause, but a similar buyer clause repeats the words in your question and gives a different code.

This ordinary lookup becomes less predictable in a long prompt. The model may fail because it must process a longer document, because the answer appears in a weak position, or because nearby text resembles the question. Each explanation points to a different diag-

nosis: compare matched prompt lengths, move the key clause, or inspect what a filter retains. Distinguishing among them requires an evaluation that holds scenario content fixed across comparisons.

We study these choices in a synthetic, gradeable version of the contract situation. Each document contains an answer-bearing vendor clause and a look-alike clause for another company. The question uses the wording of the look-alike clause, while the answer clause expresses the same relation with different words. This construction reduces the value of literal matching, following the motivation of NoLiMa (Modarressi et al., 2025).

We provide a paired evaluation of a frozen open model at laptop scale. Prior work established position and length failures in larger systems (Liu et al., 2023; Levy et al., 2024) and showed that irrelevant statements can disrupt reasoning (Shi et al., 2023). Our main comparison moves the same answer clause while preserving the question, clause wording, codes, and filler order. We then compare a deterministic keyword filter with a query-entity selector that keeps sentences naming the company in the question. This selector comparison tests whether shortening helps when the filter does or does not retain the lexical competitor.

The experiment asks three questions:

1. How does moving the same answer clause among the start, middle, and end change accuracy?
2. How much accuracy do keyword overlap and query-entity selection preserve on the same cases, and which clauses do they retain?
3. As mean prompt length grows from 1989 to 7992 tokens, how much does accuracy change on shared cases?

Within this grid, accuracy was higher when the answer appeared near the end. That placement result is diagnostic because using it requires knowing the answer location. The deployable query-entity selector instead avoided the lexical trap that keyword pruning

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

amplified, while the matched length comparison did not identify a degradation threshold.

2. Related Work

Liu et al. (2023) found a U-shaped position curve in multi-document question answering and key-value retrieval, with stronger use of information near context boundaries. We conduct a smaller clause-reordering test in which the question, filler sequence, clauses, and codes remain fixed while we insert the answer clause at a different depth. When we place it before the competing clause, the competing clause’s final sentence index shifts by one. The resulting pattern mainly favors recent information rather than forming a full U-shape, because accuracy is also low in the start condition.

Prompt length can matter even when the underlying problem is fixed. Levy et al. (2024) padded matched reasoning problems and observed degradation well before model context limits, while RULER found that usable context lengths can be shorter than advertised windows (Hsieh et al., 2024). Our prompt-length comparison includes only cases present at every tier, so the mix of scenarios remains fixed as prompts grow.

Shi et al. (2023) added irrelevant statements to arithmetic problems and found lower reasoning accuracy. Our filler bundles vary several properties together, so we use them to balance the main comparisons rather than to claim an isolated relevance effect. NoLiMa shows that synthetic retrieval tests based on literal word matches can overstate retrieval ability (Modarressi et al., 2025). We therefore paraphrase the question in the answer clause while giving the competing clause the same surface wording as the question. The tested model belongs to the instruction-tuned Qwen family (Yang et al., 2024).

3. Probe Design

Model and execution. We evaluate the frozen qwen2.5-0.5b-instruct-q8_0.gguf artifact, whose SHA checksum appears below:

ca59ca7f13d0e15a8cfa77bd17e65d24f6844b554a7b6c12e07a5f89ff76844e

We run it through llama-completion with greedy decoding, zero temperature, and a fixed seed. The complete manifest contains 180 probes, each stored as one prompt and one raw output. Existing outputs are skipped on rerun, so interrupted inference resumes by stable probe identifier.

Scenario families. For each family, seeded code generates the target company and code, a competing com-

pany and code, a question, and filler. The question asks for the renewal code needed to extend the target company’s contract. Its wording exactly matches the competing clause, whereas the answer-bearing clause calls the same value a renewal identifier. A matched literal control reverses the wording of the two clauses while retaining the same filler.

For each full-context family, we insert the answer clause near the start, middle, or end. Family generation does not use target position in any random seed. We assign the competing clause a fractional insertion index, either earlier or later in the document, that is independent of answer position and fixed across the three versions. Its final sentence index shifts by one when the answer is inserted before it. Thus, each paired position contrast changes both the answer placement and this mechanical one-sentence shift while preserving the scenario content and filler order.

Amount and filler control. The tokenizer, rather than a character estimate, determines prompt length. The three full-context tiers average 1989, 3986, and 7992 tokens. Each tier balances business filler with unrelated sentences. Business filler consists of other companies, records, and unrelated codes; unrelated filler describes places and events outside the contract domain. Filler pairs share the family, target position, question, and measured tier.

This manipulation changes several properties together. Compared with unrelated prose, the related filler contains more business entities, contract-like templates, and code-shaped answer candidates. The paired comparison therefore identifies the effect of these two generated filler conditions, not the effect of topical relevance alone.

The final tier contains fewer families because it is more expensive to run. Prompt-length estimates therefore use only the 12 cases shared by every tier rather than the unequal full-tier totals. We balance both filler bundles within each reported position and length comparison but do not interpret their contrast as a context-quality effect.

Selectors and controls. At the middle tier, a keyword filter retains sentences with at least two content words shared with the question, falling back to the highest-overlap sentences if too few qualify. A second selector extracts the company named after “contract with” in the question and keeps sentences containing that name. If extraction or matching fails, it falls back to the keyword filter. Both rules use only visible text, preserve original order, and require no answer labels.

The comparisons are restricted to the 30 cases with full context and both selector outputs. We also include 6 closed-book questions and 6 matched literal-versus-paraphrase controls.

The keyword filter retained the answer clause and lexical decoy in all 30 matched prompts. The query-entity selector retained the answer clause in 30 prompts and the decoy in 0. A paired diagnostic removes only the decoy from each keyword-pruned prompt, leaving every other selected sentence fixed. This diagnostic uses the known synthetic label; both automatic selectors do not.

4. Scoring and Comparisons

Normalized accuracy requires the target code to appear without the competing code. Exact-output compliance requires a response containing one code and no other text. A distractor capture contains the competing code without the target, and an invalid response contains both expected codes or neither. We strip only the runner’s fixed transport trailer before applying these rules.

We report descriptive results for the observed seeded scenario families. Each accuracy appears with its probe and family counts. Paired changes use the exact intersection of scenario keys and include wins, losses, and ties, where a win is a case corrected by the second condition and a loss is a newly incorrect case. The grid contains only 2 to 6 families per full-context tier, which does not support a population confidence claim. Greedy decoding is deterministic, so repeated decoding adds no sampling uncertainty.

5. Results

Answer to the first question: position. Moving the answer clause to the end produced the largest and most consistent observed gain. At the shortest tier, accuracy rose from 16.7% at the start to 75.0% at the end. The paired increase was 58.3 points across 12 cases from 6 families, with 7 wins, 0 losses, and 5 ties. At the middle tier, the same move raised accuracy by 50.0 points across 10 cases from 5 families, with 5 wins and 0 losses. The longest tier also favored the end by 50.0 points, but its 4 pairs came from only 2 families (Table 1).

Answer to the second question: lexical selection. The query-entity selector reversed the keyword filter’s failure. Full-context accuracy was 40.0%, compared with 6.7% after keyword filtering and 100.0% after query-entity selection. Against full context, query-entity se-

Table 1. Descriptive accuracy by answer-clause position. Each row compares the same cases across positions; W/L/T gives paired start-to-end outcomes.

Tokens	Start	Middle	End	Pairs	Families	W/L/T
1989	16.7	8.3	75.0	12	6	7/0/5
3986	20.0	30.0	70.0	10	5	5/0/5
7992	25.0	75.0	75.0	4	2	2/0/2

Table 2. Descriptive accuracy for matched selectors and diagnostic controls. Counts are probes/families.

Condition	Accuracy	Probes/families
Full context	40.0	30/5
Keyword pruned	6.7	30/5
Query entity	100.0	30/5
Decoy removed	90.0	30/5
Literal wording	83.3	6/6
Paraphrase wording	16.7	6/6
Closed book	0.0	6/6

lection produced 18 wins, 0 losses, and 12 ties across 30 cases from 5 families. Against keyword pruning, it gained 93.3 points, with 28 wins, 0 losses, and 2 ties over the same family count (Table 2).

Removing only the retained decoy raised keyword-pruned accuracy to 90.0%, a change of 83.3 points with 25 wins and no losses. Literal wording changed matched accuracy by 66.7 points, with 4 wins and no losses. Across all outputs, distractor capture was 40.0%, exact-output compliance was 96.7%, and invalid output was 7.2%. The diagnostic and deployable selector agree that retaining the lexical competitor, rather than deleting the answer, caused the keyword filter’s failure in this grid.

Answer to the third question: prompt length. The relationship between prompt length and accuracy remains unresolved. Accuracy was 50.0% at a mean prompt length of 1989 tokens and 58.3% at 3986 tokens. This paired change of 8.3 points comprised 2 wins, 1 losses, and 9 ties over 12 cases from 2 families. At 7992 tokens, accuracy was 58.3%, with 2 wins, 2 losses, and 8 ties relative to the middle tier (Table 3). Because the longest-tier intersection represents only 12 probes from 2 scenario families, these observations neither establish a benefit nor rule out degradation within the measured range.

6. Limitations

The findings concern one small quantized model and may not transfer to larger models or other quantization levels. Synthetic clause lookup is not natu-

Table 3. Descriptive matched prompt-length comparisons. Changes are longer minus shorter; W/L/T gives paired outcomes.

Comparison	Accuracy A	Accuracy B	Change	Pairs/families	W/L/T
Amount: short to middle	50.0	58.3	8.3	12/2	2/1/9
Amount: middle to long	58.3	58.3	0.0	12/2	2/2/8

ral contract analysis, and retrieval accuracy is not a measure of general capability. The 45-minute CPU budget capped the grid, leaving only 12 full-context probes at the longest tier and few families throughout. The estimates describe those observed families rather than a broader generator population. That sparse tier prevents a conclusion about degradation within the measured range, and the experiment cannot address longer prompts. The filler comparison confounds topical relation with semantic domain, entity and code density, candidate count, and sentence templates. A factorial design that matches those properties would be needed to identify a topical-relevance effect. Greedy decoding removes run-to-run sampling variation but does not show how sampling would change the estimates. Query-entity selection assumes that the question names one company exactly and that the relevant sentence repeats that name. It does not address aliases, pronouns, or queries without a unique named subject. The decoy-removal control also depends on an oracle label unavailable in natural documents.

7. Reproducibility

Analysis scripts generate every reported quantity from the committed probe manifest and raw model outputs. This procedure follows the principle that reports of computational evidence should expose their budget and variation (Dodge et al., 2019). The model and cached data are frozen and checksummed, and each raw output retains its stable probe identifier. Running make all regenerates results.json, the LaTeX value macros, and the paper from cached outputs; running make clean-deep && make all repeats deterministic model inference. Repository gates check frozen inputs, citation resolution, number provenance, pipeline freshness, result sanity, prose style, anonymity, and compilation.

References

- Dodge, J., Gururangan, S., Card, D., Schwartz, R., and Smith, N. A. Show your work: Improved reporting of experimental results. In Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing, 2019.
- Hsieh, C.-P., Sun, S., Kriman, S., Acharya, S., Rekesh,

D., Jia, F., Zhang, Y., and Ginsburg, B. Ruler: What’s the real context size of your long-context language models? arXiv preprint arXiv:2404.06654, 2024.

Levy, M., Jacoby, A., and Goldberg, Y. Same task, more tokens: the impact of input length on the reasoning performance of large language models. arXiv preprint arXiv:2402.14848, 2024.

Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., and Liang, P. Lost in the middle: How language models use long contexts. Transactions of the Association for Computational Linguistics, 2023.

Modarressi, A., Deilamsalehy, H., Dernoncourt, F., Bui, T., Rossi, R. A., Yoon, S., and Schütze, H. Nollima: Long-context evaluation beyond literal matching. arXiv preprint arXiv:2502.05167, 2025.

Shi, F., Chen, X., Misra, K., Scales, N., Dohan, D., Chi, E., Schärli, N., and Zhou, D. Large language models can be easily distracted by irrelevant context. In International Conference on Machine Learning, 2023.

Yang, A., Yang, B., Zhang, B., Hui, B., Zheng, B., and Yu, B. Qwen2.5 technical report. arXiv preprint arXiv:2412.15115, 2024.